Rulebook Agent: CAG vs RAG Decision Note

Nov 2, 2026 · @Subhashini

Decision

Use CAG. The six documents total about 2,300 words, roughly 4,000 tokens, which is under 2% of a typical 200K-token context window. Loading all six once and caching them is cheaper, simpler, and more accurate than building retrieval.

A single Rulebook Agent answers every question, and its questions cut across documents (an RFP requirement against a Kansas asset, a compliance rule, a staff skill, a contract clause), so seeing everything at once matters more than search would.

The document set

Six sample Word files in `Rulebook Agent\Data Sources`, mostly tables, all fictional or simplified.

Document	Words	What it holds	Revised
1 California RFP	525	Requirements R-01–R-12, dates, deliverables, evaluation, Addendum 1	Annually
6 Past bids	247	7 past bids, estimate benchmarks, lessons	Annually
4 Skills and rates	470	Staff inventory, rate card, Feb–Jul 2027 staffing plan	Quarterly
2 Kansas UX assets	359	14 reusable assets, IP status, gaps	Quarterly
5 Contract library	343	Clause playbook, fallbacks, approvers	Quarterly
3 Compliance rules	337	Rules C-01–C-12, INT-01–INT-02, checkpoints	Quarterly
Total	**2,281**	about 4,000 tokens	4 releases a year

How it scores

All five tests point to CAG today.

Test	This document set	Favours
Size	~4,000 tokens, under 2% of the window	CAG
Change rate	Quarterly at most; one rebuild every 3 months, with documents 1 and 6 riding along once a year	CAG
Question type	Cross-document: requirement → asset → rule → staff → clause	CAG
Query volume	One agent asks repeatedly over the same set, so the cache is reused	CAG
Access	Rates and labor costs (doc 4) may be restricted from some users	CAG with a separate cache per audience

Quarterly release routine: tag the version (e.g. `2027-Q1`), rebuild the cache with the annual documents 1 and 6 placed first so their part of the cache rarely changes, run a golden set of about 30 questions, then switch over and keep the last version for rollback.

When to switch to RAG

Review at each quarterly release. Move to RAG, or a hybrid, when **any one** trigger is hit:

Trigger	Today	Switch when
Corpus size	~4,000 tokens	Over ~120,000 tokens (60% of a 200K window), e.g. many full RFPs and past proposals added
Revision rate	Quarterly / annually	Documents change weekly or ad hoc, independently of each other
Number of pursuits	1 RFP (California)	Many live RFPs at once, each with its own documents
Permissions	Up to 2 audience tiers	Per-user or per-pursuit access rules

Trigger	Today	Switch when
Accuracy	To be measured	Golden-set accuracy falls below 95%

Likely first move: a hybrid. Keep documents 2–5 (the core rulebook) cached, and retrieve RFPs and past bids, the two sets that grow in number, once the archive gets large.

- [ ] Build the golden question set from the six documents
- [ ] Confirm whether doc 4 rates need a restricted audience
- [ ] Name an owner for the quarterly release